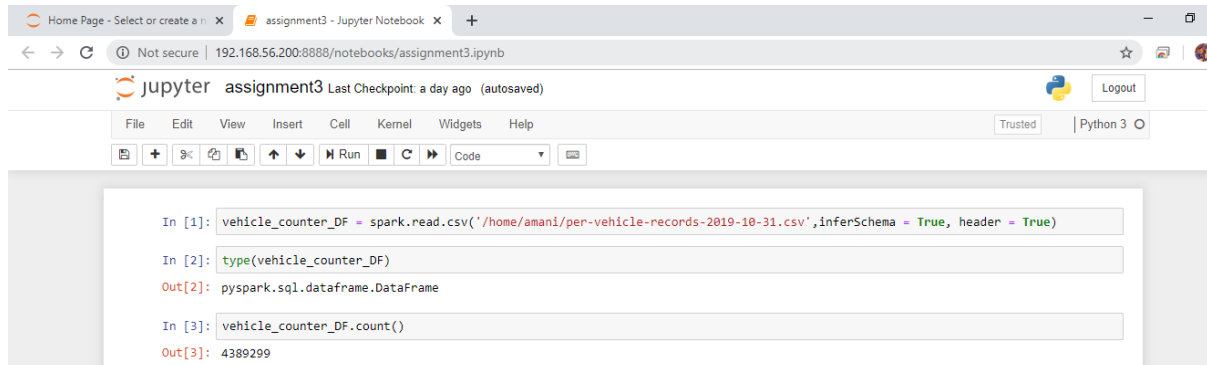


Big Data Management

Assignment 3



```
In [1]: vehicle_counter_DF = spark.read.csv('/home/amani/per-vehicle-records-2019-10-31.csv',inferSchema = True, header = True)

In [2]: type(vehicle_counter_DF)

Out[2]: pyspark.sql.dataframe.DataFrame

In [3]: vehicle_counter_DF.count()

Out[3]: 4389299
```

```
In [1]: vehicle_counter_DF = spark.read.csv('/home/amani/per-vehicle-records-2019-10-31.csv',inferSchema = True, header = True)

In [2]: type(vehicle_counter_DF)

Out[2]: pyspark.sql.dataframe.DataFrame

In [3]: vehicle_counter_DF.count()

Out[3]: 4389299
```

```
In [1]: from pyspark.streaming import StreamingContext
        from pyspark.sql import Row
        import time

In [2]: def ex1(time, rdd):
        try:
            df = spark.createDataFrame(rdd.map(\
                lambda row: Row(time=time, package=row[0], count=row[1])))
            df.write.format("org.apache.spark.sql.cassandra")\
                .options(table="question1", keyspace="streaming")\
                .save(mode="append")
        except:pass
```

```
In [4]: def save2(time, rdd):
        try:
            df = spark.createDataFrame(rdd.map(\
                lambda row: Row(time=time, word=row[0], count=row[1])))
            df.write.format("org.apache.spark.sql.cassandra")\
                .options(table="question2", keyspace="streaming")\
                .save(mode="append")
        except:pass
```

```
In [5]: def save3(time, rdd):
        try:
            df = spark.createDataFrame(rdd.map(\
                lambda row: Row(time=time, word=row[0], count=row[1])))
            df.write.format("org.apache.spark.sql.cassandra")\
                .options(table="question3", keyspace="streaming")\
                .save(mode="append")
        except:pass
```

```
In [6]: def save4(time, rdd):
        try:
            df = spark.createDataFrame(rdd.map(\
                lambda row: Row(time=time, word=row[0], count=row[1])))
            df.write.format("org.apache.spark.sql.cassandra")\
                .options(table="question4", keyspace="streaming")\
                .save(mode="append")
        except:pass
```

```
In [7]: ssc = StreamingContext(sc, 5)
```

```
amani@ubuntu:~$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.5 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cqlsh> use streaming;
cqlsh:streaming> create table question1( time text, word text, count int, PRIMAR
Y KEY(time, word));
AlreadyExists: Table 'streaming.question1' already exists
cqlsh:streaming> create table question2( time text, word text, count int, PRIMAR
Y KEY(time, word));
cqlsh:streaming> create table question3( time text, word text, count int, PRIMAR
Y KEY(time, word));
cqlsh:streaming> create table question4( time text, word text, count int, PRIMAR
Y KEY(time, word));
```

```
cqlsh:streaming> DESCRIBE TABLES;

word_counts  question1  question4  question3  question2
```

```
cqlsh:streaming> SELECT * FROM question1;

time | word | count
-----+-----+-----

(0 rows)
cqlsh:streaming> SELECT * FROM question2;

time | word | count
-----+-----+-----

(0 rows)
cqlsh:streaming> SELECT * FROM question3;

time | word | count
-----+-----+-----

(0 rows)
cqlsh:streaming> SELECT * FROM question4;

time | word | count
-----+-----+-----

(0 rows)
cqlsh:streaming> █
```

In [6]: #Show total number of counts (on each site) by vehicle class.

```
result1=spark.sql("SELECT count(class) as vehiclecount , class from vehicle_counter group by class order by count(class) desc")
result1.show()
```

vehiclecount	class
3472965	2
498505	3
216978	6
135202	5
29347	4
21224	7
14682	1
396	0

In [7]: #2. Compute the average speed (on each site) by vehicle class.

```
result2 = vehicle_counter_DF.groupBy("class")
result2.agg({'speed': 'avg'}).show()
```

class	avg(speed)
1	75.41983381010762
6	81.93572758528522
3	90.35929148153001
5	80.11806925933027
4	79.0626980611306
7	80.509602336977
2	87.99111496948547
0	81.18964646464646

In [9]: #3. Find the top 3 busiest counter sites in Ireland.

```
result3=spark.sql("SELECT cosit ,count(cosit) as cositcount from vehicle_counter group by cosit order by count(cosit) desc limit 3")
result3.show()
```

cosit	cositcount
1508	98292
1502	89498
1503	86195

In [10]: #4. Find total number of counts for HGVs.

```
result4=spark.sql("SELECT count(cosit) as HGVcount, cosit from vehicle_counter where classname in ('HGV_ART','HGV_RIG') group by cosit")
result4.show()
```

HGVcount	cosit
12031	997